


Lecture 10: Chapter 4 Part 2

The Greedy Approach
CS3310


Graph Theory

What is a minimum spanning tree of this graph?

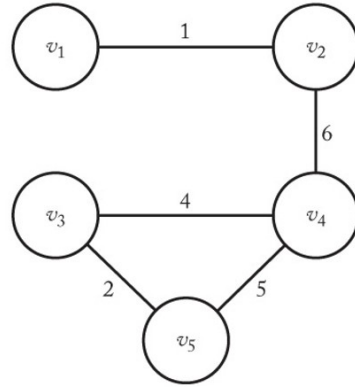


Minimum Spanning Trees

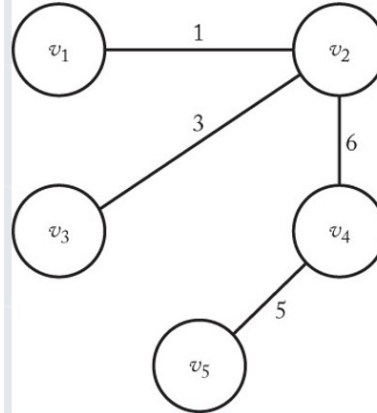
(a) A connected, weighted, undirected graph G .



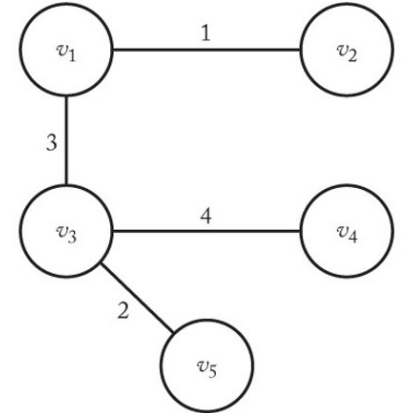
(b) If (v_4, v_5) were removed from this subgraph, the graph would remain connected.



(c) A spanning tree for G .



(d) A minimum spanning tree for G .



Kruskal's Algorithm

Like Prim's Algorithm, Kruskal's Algorithm calculates a minimum spanning tree given a weighted, connected graph $G = (V, E)$. The edges of the spanning tree are placed in F .

- Kruskal's algorithm starts by placing each vertex in V in its own disjoint set.
- For example: $V = \{1, 2, 3, 4, 5\} \rightarrow \{1\}, \{2\}, \{3\}, \{4\}, \{5\}$
- Every edge in the graph is then sorted in *nondecreasing* order of weight.
- We select the smallest remaining edge e and make sure the vertices it connects are not in the same set. If they are not, we add e to F and merge their two subsets.

Kruskal's Algorithm High-Level

```
F = {}           // initialize set of edges in spanning tree to empty

create disjoint subsets of V, one for each vertex;

sort the edges in E in nondecreasing order;


while (the instance is not solved)
    select next edge;                                     //
selection procedure
    if ( it connects 2 vertices in disjoint subsets) // feasibility check
        merge the subsets;
        add the edge to F;
    if (all the subsets are merged)                       // solution
check
    the instance is solved;
```

Kruskal's Algorithm

$V: \{v_1, v_2, v_3, v_4, v_5\}$

$E: (v_1, v_2), (v_1, v_3), (v_2, v_3), (v_2, v_4), (v_3, v_4), (v_3, v_5), (v_4, v_5)$

How do we initialize this problem?



Kruskal's Algorithm


Disjoint Sets: $\{v_1\}, \{v_2\}, \{v_3\}, \{v_4\}, \{v_5\}$

F: $\{\}$

Edges (sorted):

(v_1, v_2)	(v_3, v_5)	(v_1, v_3)	(v_2, v_3)	(v_3, v_4)	(v_4, v_5)	(v_2, v_4)
1	2	3	3	4	5	6

- Each edge is sorted by weight and each vertex is placed in its own disjoint set.
- Initialize F to the empty set.



Kruskal's Algorithm


Disjoint Sets: $\{v_1, v_2\}, \{v_3\}, \{v_4\}, \{v_5\}$

F: $\{(v_1, v_2)\}$

Edges (sorted):

(v_1, v_2)	(v_3, v_5)	(v_1, v_3)	(v_2, v_3)	(v_3, v_4)	(v_4, v_5)	(v_2, v_4)
1	2	3	3	4	5	6

- The smallest edge, (v_1, v_2) , is chosen first.
- v_1 and v_2 are in disjoint sets, so we add (v_1, v_2) to F and merge the sets v_1 and v_2 are in.



Kruskal's Algorithm


Disjoint Sets: $\{v_1, v_2\}$, $\{v_3, v_5\}$, $\{v_4\}$

F: $\{(v_1, v_2), (v_3, v_5)\}$

Edges (sorted):

(v_1, v_2)	(v_3, v_5)	(v_1, v_3)	(v_2, v_3)	(v_3, v_4)	(v_4, v_5)	(v_2, v_4)
1	2	3	3	4	5	6

- The next smallest edge, (v_3, v_5) , is chosen.
- v_3 and v_5 are in disjoint sets, so we add (v_3, v_5) to F and merge the sets v_3 and v_5 are in.



Kruskal's Algorithm


Disjoint Sets: $\{v_1, v_2, v_3, v_5\}, \{v_4\}$

F: $\{(v_1, v_2), (v_3, v_5), (v_1, v_3)\}$

Edges (sorted):

(v_1, v_2)	(v_3, v_5)	(v_1, v_3)	(v_2, v_3)	(v_3, v_4)	(v_4, v_5)	(v_2, v_4)
1	2	3	3	4	5	6

- The next smallest edge, (v_1, v_3) , is chosen.
- v_1 and v_3 are in disjoint sets, so we add (v_1, v_3) to F and merge the sets v_1 and v_3 are in.



Kruskal's Algorithm


Disjoint Sets: $\{v_1, v_2, v_3, v_5\}, \{v_4\}$

F: $\{(v_1, v_2), (v_3, v_5), (v_1, v_3)\}$

Edges (sorted):

(v_1, v_2)	(v_3, v_5)	(v_1, v_3)	(v_2, v_3)	(v_3, v_4)	(v_4, v_5)	(v_2, v_4)
1	2	3	3	4	5	6

- The next smallest edge, (v_2, v_3) , is chosen.
- v_2 and v_3 are not in disjoint sets, so we reject (v_2, v_3)



Kruskal's Algorithm

Disjoint Sets: $\{v_1, v_2, v_3, v_4, v_5\}$


F: $\{(v_1, v_2), (v_3, v_5), (v_1, v_3), (v_3, v_4)\}$

Edges (sorted): $(v_1, v_2), (v_3, v_5), (v_1, v_3), (v_2, v_3), (v_3, v_4), (v_4, v_5), (v_2, v_4)$

1 2 3 3 4 5 6

- The next smallest edge, (v_3, v_4) , is chosen.
- v_3 and v_4 are in disjoint sets, so we add (v_3, v_4) and merge the two sets

All disjoint sets are now merged, so we stop here!



Kruskal's Algorithm

- As with Prim's, Kruskal's Algorithm is easier to discuss at a high level than to implement in code.
- To make some algorithms efficient, we need a cleverly designed data structure.
- Kruskal's algorithm makes use of a **disjoint set** data structure.

Disjoint Sets for Kruskal's Algorithm

Let's say we start with a universe U of elements:

$U = \{A, B, C, D, E\}$

We can write a function `makeset` that creates a disjoint set for its argument:

```
for (each  $x \in U$ )  
    makeset( $x$ )                // make a unique set for each element of  $U$ 
```

We define a data type `set_pointer` and a function `find`. If p and q are of type

`set_pointer`:

```
 $p = \text{find}('B');$   
 $q = \text{find}('C');$ 
```


➤ p now points to the set B is in and q points to the set C is in.

Disjoint Sets for Kruskal's Algorithm

We will also define a function `merge`.

Calling `merge(p, q)` performs step b, merging two sets into one:

- (a) There are five disjoint sets. We have executed $p = \text{find}(B)$ and $q = \text{find}(C)$.



- (b) There are four disjoint sets after {B} and {C} are merged.



- (c) We have executed $p = \text{find}(B)$.




Disjoint Set Data Structure

- We will represent disjoint sets by using **inverted trees**.
- With an *inverted tree*, each nonroot points to its parent, and each root points to itself.
- Each disjoint set is represented by one tree:


(a) represents $\{A\}, \{B\}, \{C\}, \{D\}, \{E\}$

(b) represents $\{A\}, \{B, C\}, \{D\}, \{E\}$

(a) Five disjoint sets represented by inverted trees.




(b) The inverted trees after $[B]$ and $[C]$ are merged.



Disjoint Set Data Structure


One way to represent inverted trees is with an array.

- Below is array U with 10 disjoint sets.
- $U[i]$ = the parent of i
 - i.e. $U[4] = 4$ since 4 points to itself; it is the root of its tree.



Disjoint Set Data Structure


- `merge(4, 10)` places 10 in the same set as 4
 - `U[10]` is updated to 4, thus pointing node 10 to node 4.



Disjoint Set Data Structure

The following image represents:

$\{1, 5\}$, $\{2, 4, 7, 10\}$, $\{3, 6, 8, 9\}$




Disjoint Set Data Structure

Suppose we have:

$p = \text{find}(10)$

$q = \text{find}(4)$

What should `find` return?




Disjoint Set Data Structure

Suppose we have:

`p = find(10)`

`q = find(4)`


What should `find` return? The *root* of the tree the specified value is in.



Disjoint Set Data Structure

```
set_pointer find (index i)
    while (U[i] != i)
        i = U[i];
    return i;
```

This find procedure will return 2 for both `find(10)` and `find(4)`.



Disjoint Set Data Structure Definition

```
void makeset (index i)
    U[i] = i;

set_pointer find (index i)
    while (U[i] != i)
        i = U[i];
    return i;

void merge (set_pointer p, set_pointer q)
    if (p < q)
        U[q] = p;                                // p is now the parent of q
    else
        U[p] = q;                                // q is now the parent of p

bool equal (set_pointer p, set_pointer q)
    return p == q;
```

Kruskal's Algorithm

```
void kruskal(int n, int m, set_of_edges E, set_of_edges &F)
    edge e;
    Sort the m edges in E by weight in nondecreasing order;
    F = {}
    initial(n)                                // initialize n disjoint
subsets
    while (# of edges in F is less than n - 1)
        e = edge with least weight not yet considered;
        index i, j = indices of vertices connected by e;
        set_pointer p = find(i);
        set_pointer q = find(j);
        if (! equal(p, q))
            merge (p, q);
            add e to F;
```


Kruskal's Vs. Prim's

Prim's Algorithm: $T(n) \in \Theta(n^2)$

Kruskal's Algorithm: $B(m, n) \in \Theta(m \lg n)$ and $W(m, n) \in \Theta(n^2 \lg n)$


➤ m is the # of edges and n is the # of vertices

In a sparse graph with less edges, Kruskal's is $(n \lg n)$. In a dense graph with a large amount of edges, Kruskal's is $(n^2 \lg n)$

Therefore, Prim should be used with graphs with lots of edges and Kruskal's should be used with graphs that have fewer edges.

In-Class Exercise

1. Use Kruskal's Algorithm to find a minimum spanning tree for the following graph



Scheduling With Deadlines

Imagine that we have a list of jobs to be done. Each job:

- Takes 1 unit of time.
- Provides a profit if completed.
- Must be finished by a specified deadline.

If a job starts before or at its deadline, the profit is obtained. How can we schedule the jobs so that maximum profit is obtained?

Note: Not all jobs have to be scheduled. We don't consider any schedule that has a job starting after its deadline since it has the same profit as one that doesn't have that job.

➤ Such a schedule is called **impossible**.

Scheduling With Deadlines

A **Deadline** of 2 means that the job can start at time 1 *or* time 2 (there is no time 0).

- What schedules are **possible** given the following list of jobs?

Job	Deadline	Profit
1	2	30
2	1	35
3	2	25
4	1	40

Scheduling With Deadlines

A **Deadline** of 2 means that the job can start at time 1 *or* time 2 (there is no time 0).

- What schedules are **possible** given the following list of jobs?

Impossible schedules have not been listed i.e. [2, 4], [2, 1, 3] etc.

Job	Deadline	Profit
1	2	30
2	1	35
3	2	25
4	1	40

Schedule	Total Profit
[1, 3]	$30 + 25 = 55$
[2, 1]	$35 + 30 = 65$
[2, 3]	$35 + 25 = 60$
[3, 1]	$25 + 30 = 55$
[4, 1]	$40 + 30 = 70$
[4, 3]	$40 + 25 = 65$

Scheduling With Deadlines

- Schedule $[4, 1]$ is optimal. It provides a profit of 70.
 - However, considering every potential schedules takes *factorial* time!

Before we move forward, let's define a few terms:

- A sequence of jobs is a **feasible sequence** if all its jobs start by their deadlines.
 - $[4, 1]$ is feasible but $[1, 4]$ isn't; 4 has a deadline of 1 and can't start at time 2.
- A set of jobs is called a **feasible set** if there exists at least one feasible sequence for the jobs in the set.
 - $\{1, 4\}$ is a *feasible set* since $[4, 1]$ is a sequence that can be scheduled from the jobs in the set. However, $\{2, 4\}$ is not feasible since both have a deadline of 1.
- A feasible sequence with maximum total profit is called an **optimal sequence**.
- The set of jobs in that sequence is called an **optimal set of jobs**.

Scheduling With Deadlines

High-Level pseudocode:

```
sort the jobs in nonincreasing order by profit;
```

```
S = {};
```

```
while (the instance is not solved)
    select next job;
    if (S is feasible with this job added)
        add this job to S;
    if (there are no more jobs)
        the instance is solved;
```

Scheduling With Deadlines

Job	Deadline	Profit
1	3	40
2	1	35
3	1	30
4	3	25
5	1	20
6	3	15
7	2	10

1. S is set to $\{\}$
2. S is set to $\{1\}$, because $[1]$ is feasible
3. S is set to $\{1, 2\}$ because $[2, 1]$ is feasible
4. $\{1, 2, 3\}$ is rejected because there is no feasible sequence for this set
5. S is set to $\{1, 2, 4\}$ because $[2, 1, 4]$ is feasible
6. $\{1, 2, 4, 5\}$ is rejected.
7. $\{1, 2, 4, 6\}$ is rejected
8. $\{1, 2, 4, 7\}$ is rejected

The final value of S is $\{1, 2, 4\}$, and a feasible sequence for S is $[2, 1, 4]$ (we could also use $[2, 4, 1]$)

Total Profit: 100

Scheduling With Deadlines

When we add a job to the set S , we need to determine if it is still feasible. This is easy for a human on small sets, but takes factorial time for a computer.

What is an efficient way to determine this?

Scheduling With Deadlines

When we add a job to the set S , we need to determine if it is still feasible. This is easy for a human on small sets, but takes factorial time for a computer.

What is an efficient way to determine this?

- Sort the jobs in S in nondecreasing order by deadline.

Suppose we have the following set: $\{1, 2, 4, 7\}$

- Job 1 has a deadline of 3, job 2 a deadline of 1, job 4 a deadline of 3, and job 7 a deadline of 2.
- We sort the jobs in nondecreasing order by deadline: $[2, 7, 1, 4]$
 - Job 4 starts at time 4, but its deadline is 3...therefore, the set is not feasible.

Scheduling With Deadlines

Problem: determine the schedule with maximum total profit give that each job has a profit that will be obtained only if the job is scheduled by its deadline.

Inputs: n , the number of jobs. Array of integers `deadline`, indexed from 1 to n . `deadline[i]` is the deadline for the i th job. The jobs should be sorted in nonincreasing order by their profits.

Outputs: an optimal sequence `finalSequence` for the jobs.

Scheduling With Deadlines

```
void schedule (int n, const int deadline[], sequence_of_integers &finalSequence)
    index i;
    sequence_of_integers temp;

    finalSequence = [1];           // initialize finalSequence to contain the first
job
    for (i = 2; i <= n; i++)
        temp = finalSequence with job i added;
    sort temp in nondecreasing order by deadline[i];
    if (temp is feasible)
        finalSequence = temp;
```

Scheduling With Deadlines

Job	Deadline	Profit
1	3	40
2	1	35
3	1	30
4	3	25
5	1	20
6	3	15
7	2	10

1. `finalSequence` is set to [1]
2. `temp` is set to [1, 2] and sorted to [2, 1]. It is feasible.
`finalSequence` is set to [2, 1] since `temp` is feasible.
3. `temp` is set to [2, 1, 3] and sorted to [2, 3, 1]. It is not feasible and is rejected.
4. `temp` is set to [2, 1, 4] and is already sorted.
`finalSequence` is set to [2, 1, 4] because `temp` is feasible.
5. `temp` is set to [2, 1, 4, 5], sorted to [2, 5, 1, 4]. Rejected.
6. `temp` is set to [2, 1, 4, 6], sorted to [2, 1, 6, 4]. Rejected.
7. `temp` is set to [2, 1, 4, 7], sorted to [2, 7, 1, 4]. Rejected.

The final value of `finalSequence` is [2, 1, 4].

In-Class Exercise

1. Consider the following jobs, deadlines, and profits. Use the Scheduling with Deadlines algorithm to maximize the total profit

Job	Deadline	Profit
1	2	40
2	4	15
3	3	60
4	2	20
5	3	10
6	1	45
7	1	55